

Homework 2

- Yucheng Zhang

0. Machine

- MacBook Pro 2017
- Compiler version: `g++ (Homebrew GCC 8.2.0) 8.2.0`.
- CPU: `Intel(R) Core(TM) i7-7700HQ CPU @ 2.80GHz`.
 - Details: <https://ark.intel.com/products/97185/Intel-Core-i7-7700HQ-Processor-6M-Cache-up-to-3-80-GHz>.
 - Cores: `4`, Threads: `8`.
 - Max turbo frequency: `3.80 GHz`.
 - Operations per cycle: `16 DP FLOPs/cycle` for Intel Kaby Lake as found here <https://stackoverflow.com/a/15657772>.
 - Max flop rate: `Cores * Max turbo frequency * Operations per cycle = 243.2 Gflop/s`.
 - Max memory bandwidth: `37.5 GB/s` for 2 channels.

1. Finding Memory bugs.

1. Index exceeds the array range; `free()` should be used for memory allocated by `malloc`.
2. Variable is used before initialized, while valgrind does not complain about this.

2. Optimizing matrix-matrix multiplication.

- Try different loop arrangements.
 - The outputs are shown in [files/2-order-jpi.txt](#), [files/2-order-jjp.txt](#) and [files/2-order-ipj.txt](#).
 - We can see that the performance of order jpi is the best. This is because the matrices are stored in column major order. For order jpi, `double A_ip = a[i + p * m]; double B_pj = b[p + j * k]; double C_ij = c[i + j * m];` are all read in continuous memory location, which can save a lot of time.
 - Similarly, we get the worst performance for order ipj, which would be the best if the matrices were stored in row major order.
- Implement a one level blocking scheme by using `BLOCK_SIZE` macro as the block size.
- Experiment with different values for `BLOCK_SIZE`.
 - From the table below, we can see that we get better performance around `BLOCK_SIZE = 64`.

BLOCK_SIZE	Gflop/s (Average)	GB/s (Average)
4	6.089491	97.431855
8	3.242940	51.887047
16	2.426508	38.824121
32	15.953180	255.250880
64	19.541914	312.670619
128	16.749192	267.987077
256	14.572623	233.161976

- Parallelize your matrix-matrix multiplication code using OpenMP.
 - I parallelize the code on the for loop over blocks in C.
 - One thing to notice is how the cache is shared by all the threads. The optimal BLOCK_SIZE may be different when we use different number of threads. For example, when I use OpenMP with more than one thread, BLOCK_SIZE = 32 works better than BLOCK_SIZE = 64, which is optimal in the serial case.
- What percentage of the peak FLOP-rate do you achieve with your code?
 - I can achieve up to 22.6 % of the peak FLOP-rate.

3. Finding OpenMP bugs.

2. reduction should be used for simple sum. int and float may not be large enough in some case, which might depend on the machine.
3. For a function which may not be executed by all threads, #pragma omp barrier inside may cause the program to get stuck.
4. private stack size is not very large.
5. lock may cause the program to get stuck if not used properly.
6. We may use global variables in order to be shared easily.

4. OpenMP version of 2D Jacobi/Gauss-Seidel smoothing.

- The following tables show the timings for different values of N and different numbers of threads. Number of iterations is 100.
- Jacobi method

N_thread	N=100	N=1000	N=10,000	N=20,000
1	0.002660 s	0.209642 s	19.760169 s	81.583017 s
2	0.004587 s	0.163366 s	13.533498 s	54.148035 s
4	0.005771 s	0.136099 s	12.947494 s	53.681173 s
8	0.007180 s	0.143475 s	13.939527 s	69.407606 s

- Gauss-Seidel method

N_thread	N=100	N=1000	N=10,000	N=20,000
1	0.001833 s	0.204328 s	23.346515 s	94.526992 s
2	0.006006 s	0.186694 s	19.875444 s	78.298179 s
4	0.009349 s	0.177580 s	19.546787 s	79.597596 s
8	0.011964 s	0.199010 s	20.717127 s	81.617238 s

- We can see that for large N, 2 threads work better than 1 thread, but more threads (4 and 8) don't really perform better. I think it's related to the memory bandwidth and also the fact that I'm using other softwares during the timing.
- I also test the code on another machine, which has Intel(R) Xeon(R) CPU E5-2650 v4 @ 2.20GHz, with 24 cores and 1 thread per core. Timings are summarized below, with 100 iterations and N=20,000.
- Jacobi method

N_thread	1	2	4	8	16	24
Time (s)	137.929256	66.865910	38.367468	25.794236	21.755469	21.025227

- Gauss-Seidel method

N_thread	1	2	4	8	16	24
Time (s)	172.466722	81.096346 s	43.46177	26.440732	21.187021	21.504986

- We see that on this more stable machine, we can reduce the time to half if we double the number of threads until N_thread = 8. For more threads, the performance may be restricted by the cache size and bandwidth.